MODEL EVALUATION · DOCUMENT AI

Invoice Field Extraction — Model Benchmark

September 3, 2026 | 14 model configurations | 29,308 documents scored | 15 fields per invoice | $128.49 total API spend

Key results

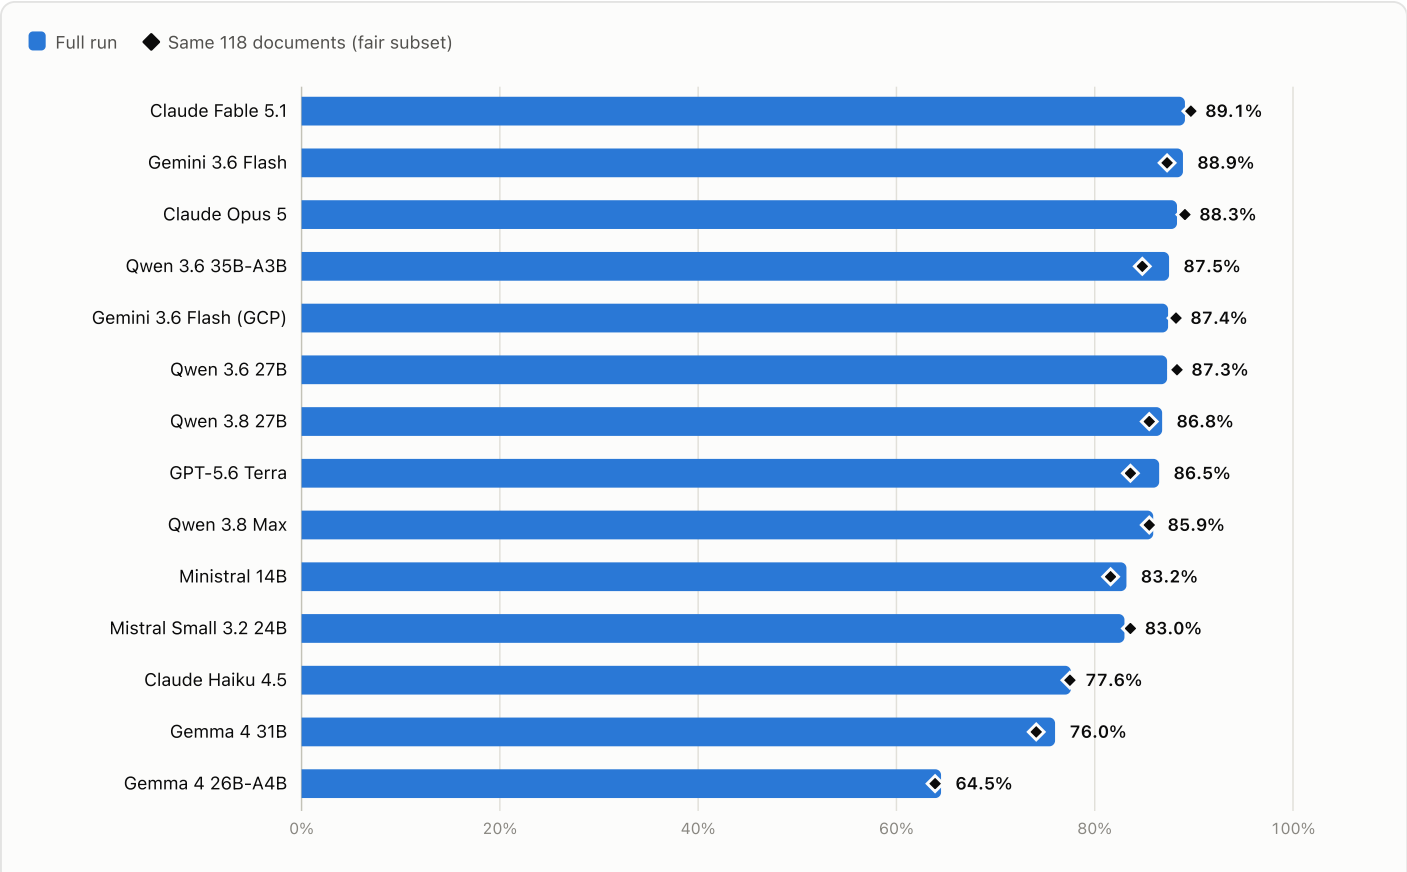
Each model reads invoice PDFs as page images and returns 15 structured fields (vendor, customer, totals, dates, references). The headline metric is **end-to-end score**: the share of all fields present in a document that the model extracted correctly.

TOP SCORE	BEST VALUE	FASTEST RESPONSE	PREMIUM CEILING
89.1%	**87.3%**	**6.9 s**	**92.2%**
Claude Fable 5.1 · $77.67 / 1k docs · 207 docs	Qwen 3.6 27B · $2.20 / 1k docs, open-weight	GPT-5.6 Terra · 86.5% E2E · $8.27 / 1k docs	Claude Opus 5 accuracy — at $44.19 / 1k docs

Bottom line: Claude Fable 5.1 posts the highest end-to-end score (89.1%) and holds the top spot on the fair-subset re-score too — but that's on only 207 documents at $77.67/1k, a thin and expensive sample next to the rest of the field. **Gemini 3.6 Flash** remains the strongest production pick: a large, cheap run (2,323 docs, $10.04/1k) that trails the top score by just 0.2 points. **Qwen 3.6 27B** is still the best open-weight value at roughly a fifth of Gemini's cost with only 1.6 points less accuracy — its newer sibling **Qwen 3.8 27B** (86.8%) doesn't improve on it despite the version bump. **Mistral Small 3.2 24B** is the cheapest credible option tested at $0.39/1k for 83.0%. **GPT-5.6 Terra** is still the fastest model tested, and **Claude Opus 5** still gives the most accurate individual answers (92.2%), but at 4–20× the unit cost of the mid-tier options it only makes sense for high-value documents.

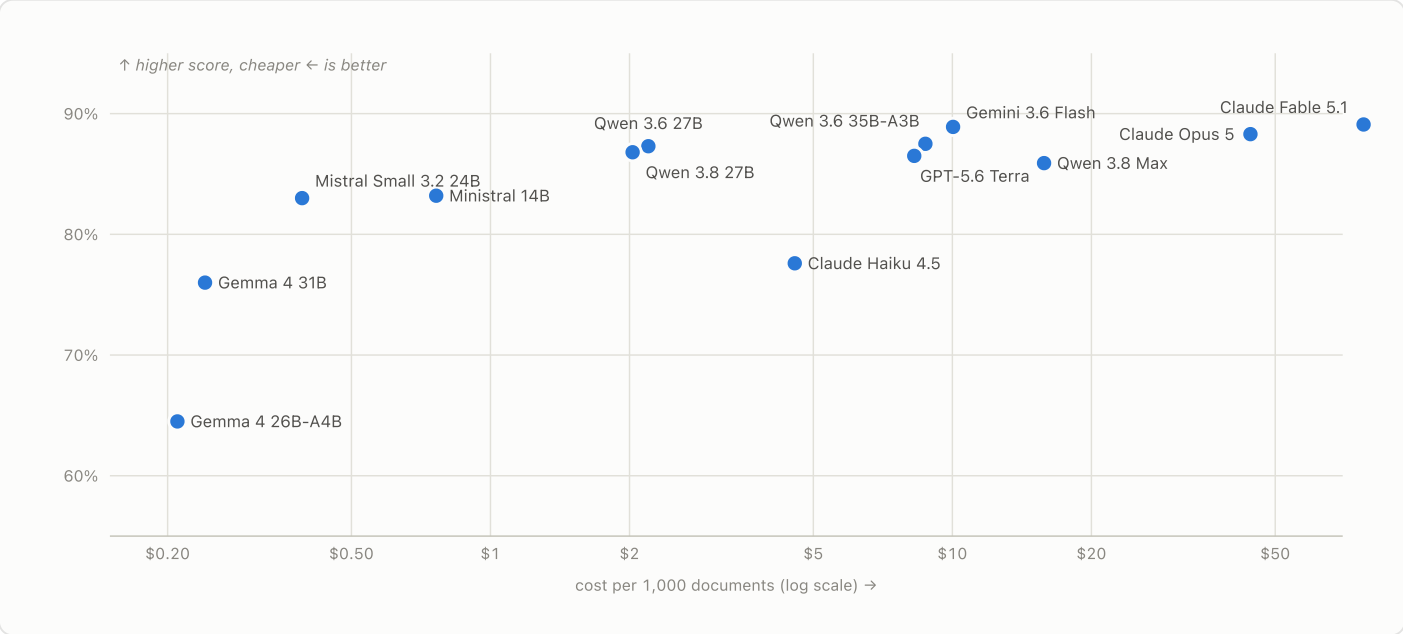
Leaderboard — end-to-end extraction score

Bars show each model on its full run. Because models processed different document counts, diamonds re-score every model on the identical 118 documents all fourteen completed — the ranking mostly holds, so the full-run numbers are trustworthy.



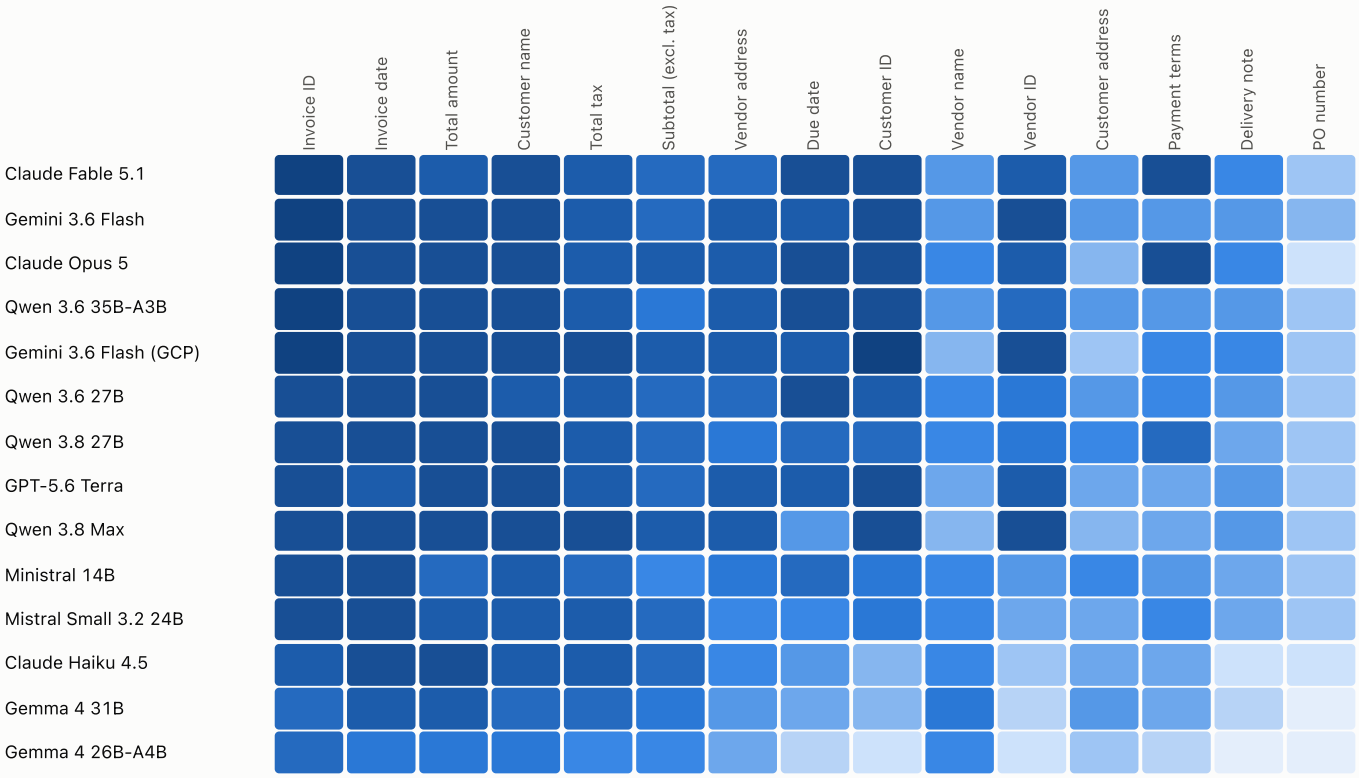
Cost vs. quality

What one thousand documents cost against what you get back. Up and to the left is better; the cost axis is logarithmic. Gemini 3.6 Flash (GCP deployment) is omitted — it is billed through the GCP contract, so no per-call cost is reported.



Where extraction breaks down, field by field

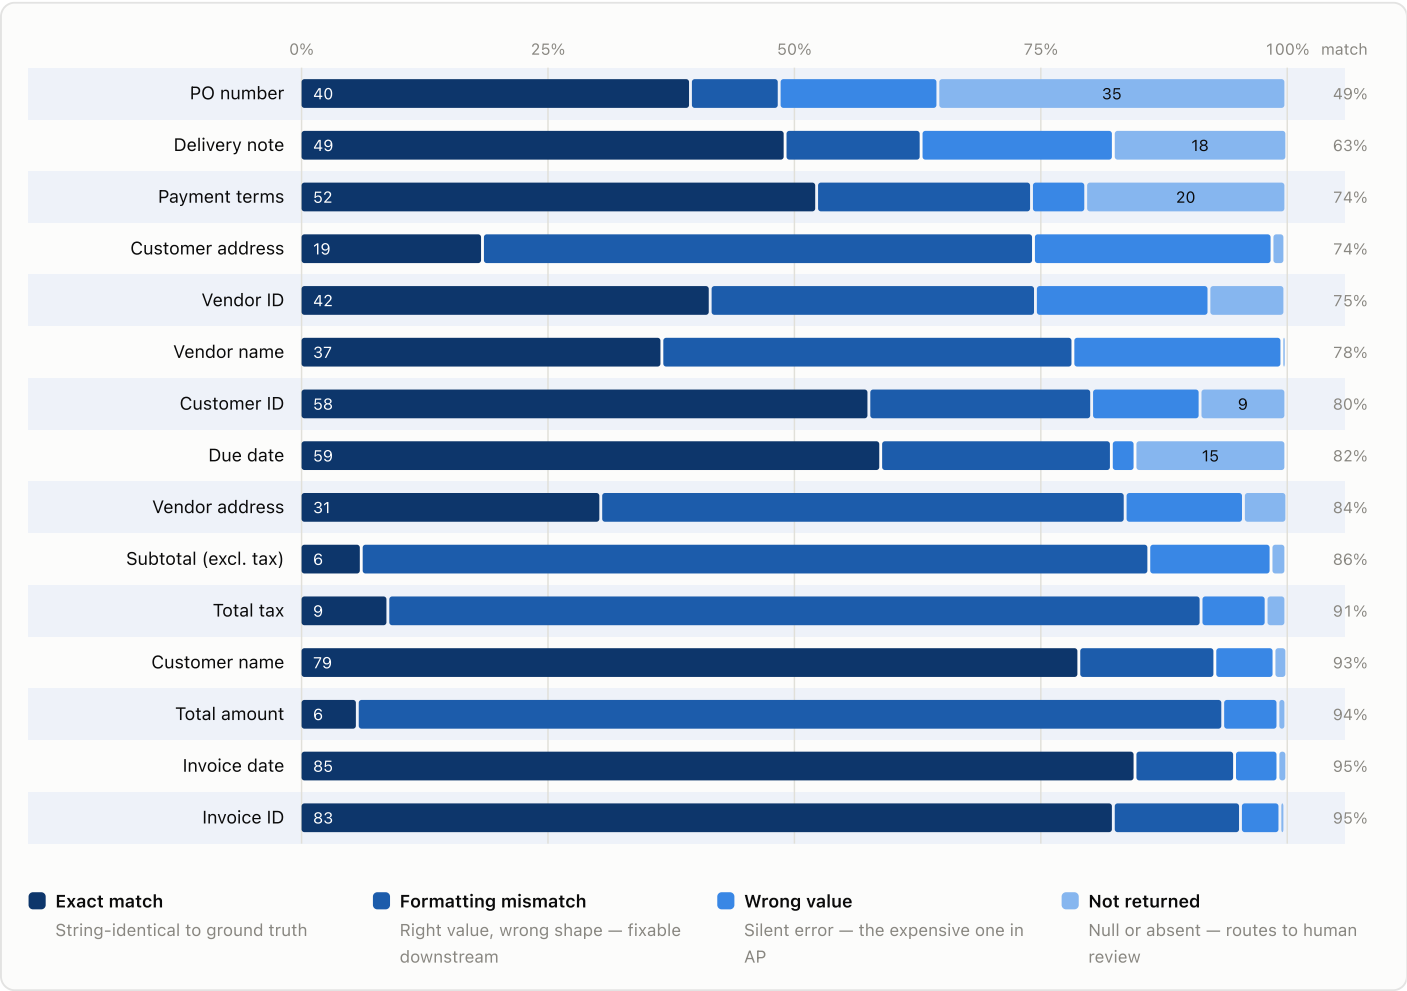
End-to-end score per field. Header fields (invoice ID, dates, totals) are largely solved; the losses concentrate in **PO number, delivery note and payment terms** — mostly fields the models fail to return at all — and in vendor-name formatting mismatches. Hover any cell for exact numbers.



▶ Field-level scores as a table

What kind of failure is it?

Every ground-truth field across all model runs combined — ≈395,000 comparisons — split by outcome. The totals are almost never string-identical (ground truth stores raw numbers, models copy the printed format) but overwhelmingly correct after normalisation. Vendor name and the addresses are mostly a normalisation problem; **PO number, delivery note and payment terms are a coverage problem** — the model returns nothing at all.



Full comparison

MODEL	END-TO-END	E2E (118 DOCS)	COVERAGE	ACCURACY	COST / 1K DOCS	TOKENS / DOC	LATENCY	DOCS RUN
Claude Fable 5.1	**89.1%**	**89.7%**	**97.0%**	91.9%	$77.67	6,417	11.1 s	207
Gemini 3.6 Flash	88.9%	87.3%	**97.0%**	91.7%	$10.04	3,164	11.0 s	2,323
Claude Opus 5	88.3%	89.1%	95.7%	**92.2%**	$44.19	7,473	16.7 s	234
Qwen 3.6 35B-A3B	87.5%	84.8%	96.8%	90.3%	$8.74	11,241	82.9 s	2,361
Gemini 3.6 Flash (GCP)	87.4%	88.2%	96.9%	90.2%	n/a	5,113	12.3 s	488
Qwen 3.6 27B	87.3%	88.3%	96.8%	90.2%	$2.20	4,575	42.5 s	2,380
Qwen 3.8 27B	86.8%	85.5%	96.3%	90.1%	$2.03	4,204	26.7 s	3,508
GPT-5.6 Terra	86.5%	83.6%	96.0%	90.0%	$8.27	5,543	6.9 s	1,815
Qwen 3.8 Max	85.9%	85.5%	95.0%	90.4%	$15.79	5,866	22.9 s	950
Ministral 14B	83.2%	81.6%	95.2%	87.3%	$0.76	4,127	13.2 s	4,973
Mistral Small 3.2 24B	83.0%	83.6%	94.0%	88.3%	$0.39	4,723	11.4 s	1,300
Claude Haiku 4.5	77.6%	77.5%	90.0%	86.3%	$4.56	3,415	509 s*	2,182
Gemma 4 31B	76.0%	74.1%	90.7%	83.8%	$0.24	1,067	18.3 s	4,158
Gemma 4 26B-A4B	64.5%	63.9%	78.7%	82.0%	**$0.21**	1,078	37.3 s	2,429

Coverage = share of fields present in the document that the model returned. Accuracy = share of returned values that match ground truth. End-to-end = correct ÷ present. Best column value in blue. * Claude Haiku 4.5 latency is inflated by rate-limit retries during its run and is not comparable.

Method & caveats

- **Setup.** Invoice PDFs rendered locally to page images (up to 4 pages per document, 150 dpi) and sent to each model with the same extraction prompt and JSON schema, temperature 0, via OpenRouter — plus one Gemini 3.6 Flash deployment on GCP Vertex for comparison. Reasoning modes were capped at minimum effort so models are compared on extraction, not thinking-token budgets. Two endpoints (Qwen 3.8 Max, Claude Fable 5.1) make reasoning mandatory and reject it disabled outright — those two ran with reasoning forced on at minimum effort, everything else with reasoning off.
- **Ground truth.** Model output is compared field-by-field against a reviewed ground-truth set, with tolerant matching for dates and number formats.
- **Document category is excluded** from all scores. Every model scored ~0% on it because the ground truth uses internal labels (e.g. “Expense”) while models return the document’s printed type — a taxonomy mismatch to fix in the ground truth, not a model failure.
- **Unequal run sizes.** Runs range from 207 documents (Claude Fable 5.1) to 4,973 (Ministral 14B) because of per-model cost caps. The fair-subset re-score on the 118 shared documents mostly confirms the ordering — Claude Fable 5.1 and Claude Opus 5 hold the top two spots either way — though Gemini 3.6 Flash drops from 2nd to 5th on the smaller sample while Qwen 3.6 27B and the GCP Gemini deployment move up into 3rd and 4th.
- **Vendor name scores understate the models.** A large share of vendor-name misses are formatting variants of the correct name; tightening ground-truth normalization would lift every model roughly equally.
- **Failure types** (in “What kind of failure is it?”): *exact* = string-identical to ground truth; *formatting mismatch* = same value after normalisation (date formats, thousand separators, accents, punctuation) — these count as correct in all scores; *wrong value* = returned but doesn’t match; *not returned* = null or absent. Pooled across all fourteen models’ runs.
- **Costs** are actual billed OpenRouter charges at list prices, not estimates.
- **Meta Muse Spark 1.2 was not benchmarked.** OpenRouter gates that endpoint behind an account-level 18+ age attestation that wasn't completed before this report was built, so every call returned HTTP 403 with zero data collected — it is omitted here rather than shown with a misleading 0% score.